



Using SVD to Predict Movie Ratings

Serdar Sali

sali@soe.ucsc.edu



Collaborative filtering



- Similar tastes in the past : similar tastes in the future

Customers Who Bought This Item Also Bought



[HP Pavilion TX1420US 12.1\"/>
Laptop \(AMD Turion 64 X2
Dual-Core...](#)

★★★★☆ (11) \$1,169.99



[HP Pavilion DV6775US 15.4\"/>
Entertainment Laptop \(Intel
Core 2...](#)

★★★★☆ (18) \$1,099.99



[HP Pavilion DV2740SE 14.1\"/>
Entertainment Laptop \(AMD
Turion 64...](#)

★★★★☆ (12) \$1,179.83



[HP Pavilion TX1320US 12.1\"/>
Entertainment Notebook PC
\(AMD Turi...](#)


★★★★☆ (35) \$1,199.99



[HP Artist Edition Messenger
Case](#)


\$49.99

Our domain: Movie Ratings




- We would like to predict how a user would rate a given movie
- A multi-label classification problem
 - Ratings : 1 to 5
 - Only available data : User – movie – rating triplets

Problem Formulation




- R : $u \times m$ matrix of ratings
 - u : number of users
 - m : number of movies
- Ideally, we would have:
 - A set of features for each movie: f_j
 - A set of preference multipliers for each user : call p_i
 - Then rating of user i for movie j becomes $r_{ij} = p_i f_j^T$

Problems




- Feature list for movies are hard to obtain
 - Task is inherently subjective and difficult, classification is hard
 - Dependence on external data resources
 - Tremendous effort required to clean-up data
- Notice that R already contains this data, but it is lumped together in a sum.
 - Can we retrieve it somehow?

Singular Value Decomposition




- SVD states that every $m \times n$ matrix A can be written as $A = USV^T$ where
 - U is an $m \times m$ orthogonal matrix,
 - S is an $m \times n$ diagonal matrix with singular values of A along the diagonal,
 - V is an $n \times n$ orthogonal matrix.


Full vs Reduced SVD



- Since S is diagonal, we can obtain a more compact representation:




SVD for Matrix Approximation




- Instead of using all the singular values of S , use only the most significant r
- Compute a rank- r approximation A' to A such that $A' = U'S'V'^T$ where U' is $m \times r$, S' is $r \times r$, and V' is $n \times r$
- This approximation minimizes the Frobenius form: $\|A - A'\|_F = \sqrt{\sum (a_{ij} - a'_{ij})^2}$

SVD for Movie Rating Prediction




- Given a matrix of ratings R , we want to compute an approximate matrix R_{app} such that RMSE is minimized.
- But $\text{RMSE} = \|R - R_{\text{app}}\|_F$
- So, SVD is a perfect fit to our problem

SVD for Movie Rating Prediction



- Recall R_{uxm} : ratings matrix
- Compute an SVD for R and just lump the singular value matrix in the sum:
 - $R = P_{uxf} F_{mxf}^T$
 - P : Preference matrix for f features for u users
 - F : f -features matrix for m movies

But...




- SVD is not defined for sparse matrices
 - Netflix data: 8.5B possible entries, 8.4B empty
- Fill in with averages, some clever combinations
 - Perturbs the data too much
 - And even if we fill in the missing values...
- Computing SVD for large matrices is computationally very expensive

Incremental SVD Method




- Devised by Simon Funk
- Only consider existing values
- Do a gradient descent to minimize the error:
 - $E = (R - R_{app})_{ij}^2$
 - Take the derivative wrt p_{ij} and f_{jk} , and the updates become
 - $p_{ik}^{(t+1)} = p_{ik}^{(t)} + \text{learning_rate} * (R - R_{app})_{ij} * f_{jk}^{(t)}$
 - $f_{jk}^{(t+1)} = f_{jk}^{(t)} + \text{learning_rate} * (R - R_{app})_{ij} * p_{ik}^{(t)}$

Implementation




- Online update
- Set each feature & each multiplier to 0.1
- Train the most significant feature first, and then the second, etc.
- Parameters:
 - Number of features
 - Learning rate
 - Regularization : very simple : $-K^*(\text{update target})$
 - Different starting values

Experiments

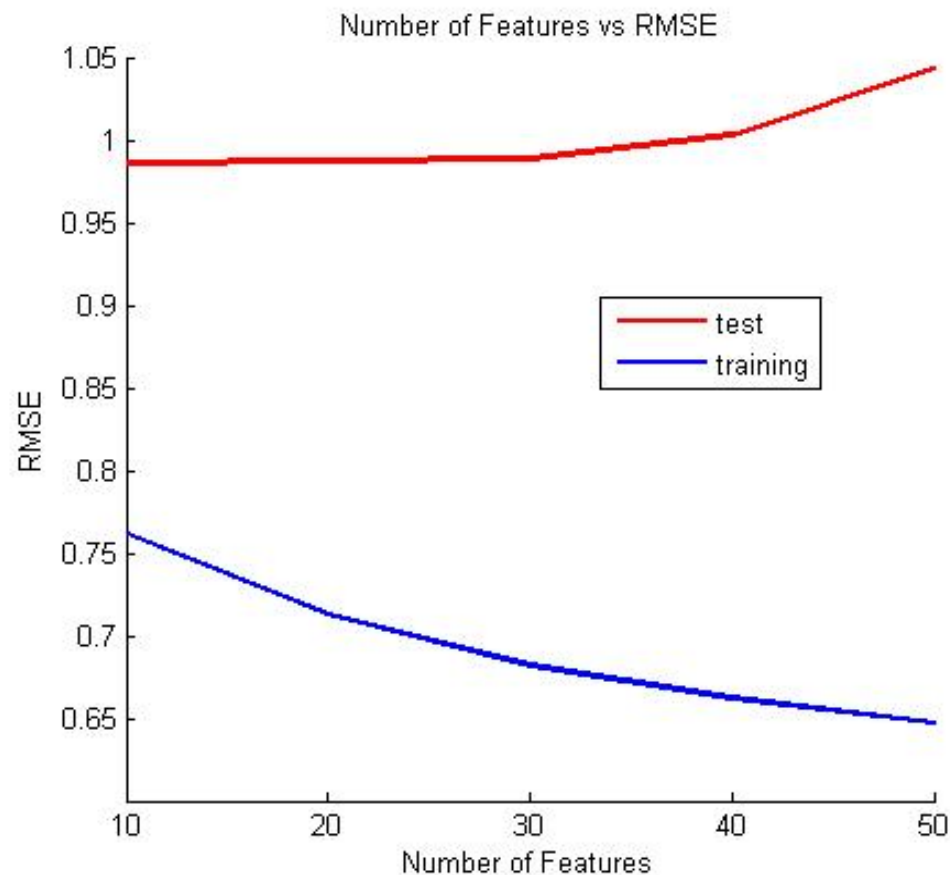


- Smaller dataset
 - 2000 movies, 480189 users
 - 10314269 ratings
 - Just blind downsampling
- Test: predict 3000 ratings
- Does not perform as well as it does on whole dataset

Experiments




- Number of Features

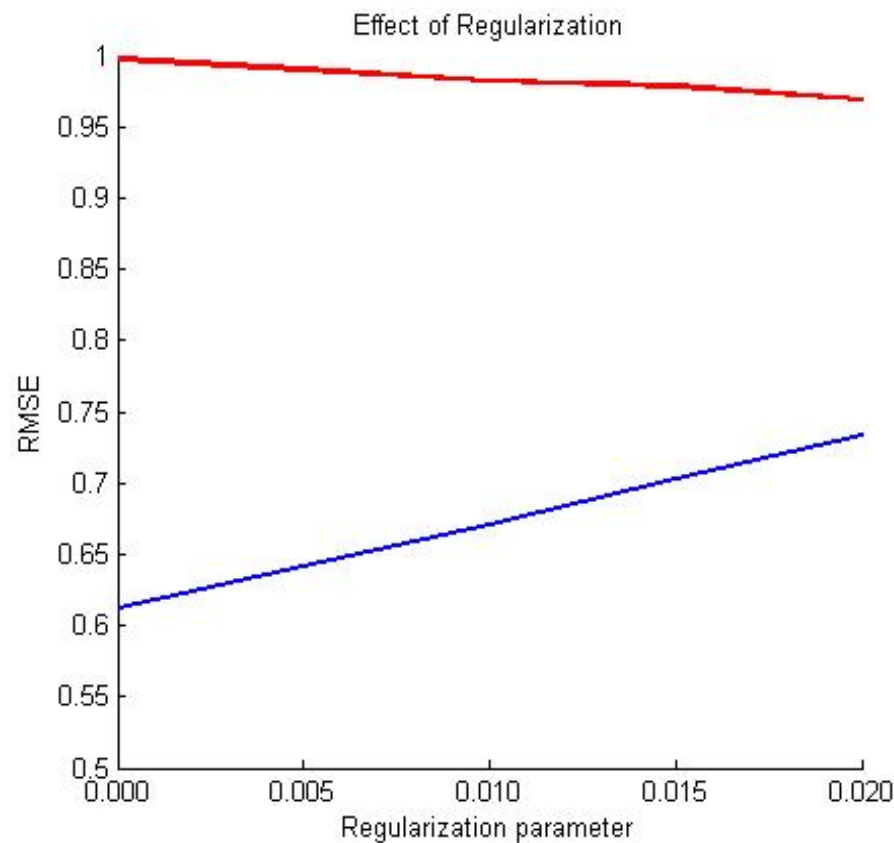


#feats	Training	Test
10	0.7623	0.9858
20	0.7128	0.9869
30	0.6820	0.9882
40	0.6619	1.0028
50	0.6468	1.0438
Combined : 0.9895		

Experiments




- Regularization




Reg Rate	Training	Test
0.0000	0.6118	0.9973
0.005	0.6418	0.9899
0.010	0.6705	0.9823
0.015	0.7027	0.9787
0.020	0.7333	0.9686
Combined : 0.9741		

Experiments




- Different learning rates



Lrate	Training	Test
0.0005	0.8180	0.9755
0.0010	0.7666	0.9765
0.0015	0.7464	0.9843
0.0020	0.7365	0.9904
0.0025	0.7314	0.9955
Combined : 0.9756		

Experiments




- Different starting values

	Base = 1	Base = Average	Base = Average + Offset
Training	0.7651	0.7638	0.7638
Test data	1.0144	1.0156	1.0158


- Combined : 0.9741

Conclusion



- Overfitting is a problem
 - Especially with my smaller dataset, models tend to overfit the training data very easily
- Even a blind combination of results give surprisingly good results
 - This implies that different models work good for different cases
 - A combination of different models is the way to go

Future Work



- Many parameters to adjust
- More clever downsampling of the dataset
- Use the computed features as input to another algorithm?
 - May help fine-tune the results at least
- Different regularization methods

Questions?

